

```
!pip install supervision==0.19.0
```

```
Requirement already satisfied: supervision==0.19.0 in /usr/local/lib/python3.10/dist-packages (0.19.0)
Requirement already satisfied: defusedxml<0.8.0,>=0.7.1 in /usr/local/lib/python3.10/dist-packages (from supervision==0.19.0) (0.7.1)
Requirement already satisfied: matplotlib>=3.6.0 in /usr/local/lib/python3.10/dist-packages (from supervision==0.19.0) (3.7.1)
Requirement already satisfied: numpy>=1.21.2 in /usr/local/lib/python3.10/dist-packages (from supervision==0.19.0) (1.25.2)
Requirement already satisfied: opencv-python-headless>=4.5.5.64 in /usr/local/lib/python3.10/dist-packages (from supervision==0.19.0) (4.5.5.64)
Requirement already satisfied: pillow>=9.4 in /usr/local/lib/python3.10/dist-packages (from supervision==0.19.0) (9.4.0)
Requirement already satisfied: pyyaml>=5.3 in /usr/local/lib/python3.10/dist-packages (from supervision==0.19.0) (6.0.1)
Requirement already satisfied: scipy<2.0.0,>=1.10.0 in /usr/local/lib/python3.10/dist-packages (from supervision==0.19.0) (1.11.4)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=3.6.0->supervision==0.19.0) (1.0.1)
Requirement already satisfied: cycycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=3.6.0->supervision==0.19.0) (0.10.0)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=3.6.0->supervision==0.19.0) (4.22.0)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=3.6.0->supervision==0.19.0) (1.0.1)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=3.6.0->supervision==0.19.0) (23.1)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=3.6.0->supervision==0.19.0) (3.1.0)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=3.6.0->supervision==0.19.0) (2.8.2)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.7->matplotlib>=3.6.0->supervision==0.19.0) (1.16.0)
```

```
!pip install ultralytics
```

```
Requirement already satisfied: ultralytics in /usr/local/lib/python3.10/dist-packages (8.1.37)
Requirement already satisfied: matplotlib>=3.3.0 in /usr/local/lib/python3.10/dist-packages (from ultralytics) (3.7.1)
Requirement already satisfied: opencv-python>=4.6.0 in /usr/local/lib/python3.10/dist-packages (from ultralytics) (4.8.0.76)
Requirement already satisfied: pillow>=7.1.2 in /usr/local/lib/python3.10/dist-packages (from ultralytics) (9.4.0)
Requirement already satisfied: pyyaml>=5.3.1 in /usr/local/lib/python3.10/dist-packages (from ultralytics) (6.0.1)
Requirement already satisfied: requests>=2.23.0 in /usr/local/lib/python3.10/dist-packages (from ultralytics) (2.31.0)
Requirement already satisfied: scipy>=1.4.1 in /usr/local/lib/python3.10/dist-packages (from ultralytics) (1.11.4)
Requirement already satisfied: torch>=1.8.0 in /usr/local/lib/python3.10/dist-packages (from ultralytics) (2.2.1+cu121)
Requirement already satisfied: torchvision>=0.9.0 in /usr/local/lib/python3.10/dist-packages (from ultralytics) (0.17.1+cu121)
Requirement already satisfied: tqdm>=4.64.0 in /usr/local/lib/python3.10/dist-packages (from ultralytics) (4.66.2)
Requirement already satisfied: psutil in /usr/local/lib/python3.10/dist-packages (from ultralytics) (5.9.5)
Requirement already satisfied: py-cpuinfo in /usr/local/lib/python3.10/dist-packages (from ultralytics) (9.0.0)
Requirement already satisfied: thop>=0.1.1 in /usr/local/lib/python3.10/dist-packages (from ultralytics) (0.1.1.post2209072238)
Requirement already satisfied: pandas>=1.1.4 in /usr/local/lib/python3.10/dist-packages (from ultralytics) (1.5.3)
Requirement already satisfied: seaborn>=0.11.0 in /usr/local/lib/python3.10/dist-packages (from ultralytics) (0.13.1)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=3.3.0->ultralytics) (1.0.1)
Requirement already satisfied: cycycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=3.3.0->ultralytics) (0.10.0)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=3.3.0->ultralytics) (4.22.0)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=3.3.0->ultralytics) (1.0.1)
Requirement already satisfied: numpy>=1.20 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=3.3.0->ultralytics) (1.25.2)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=3.3.0->ultralytics) (23.1)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=3.3.0->ultralytics) (3.1.0)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/dist-packages (from matplotlib>=3.3.0->ultralytics) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas>=1.1.4->ultralytics) (2023.4)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests>=2.23.0->ultralytics) (3.3.2)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests>=2.23.0->ultralytics) (3.6)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests>=2.23.0->ultralytics) (2.0.7)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests>=2.23.0->ultralytics) (2023.7.22)
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from torch>=1.8.0->ultralytics) (3.13.1)
Requirement already satisfied: typing-extensions>=4.8.0 in /usr/local/lib/python3.10/dist-packages (from torch>=1.8.0->ultralytics) (4.8.0)
Requirement already satisfied: sympy in /usr/local/lib/python3.10/dist-packages (from torch>=1.8.0->ultralytics) (1.12)
Requirement already satisfied: networkx in /usr/local/lib/python3.10/dist-packages (from torch>=1.8.0->ultralytics) (3.2.1)
Requirement already satisfied: jinja2 in /usr/local/lib/python3.10/dist-packages (from torch>=1.8.0->ultralytics) (3.1.3)
Requirement already satisfied: fsspec in /usr/local/lib/python3.10/dist-packages (from torch>=1.8.0->ultralytics) (2023.6.0)
Requirement already satisfied: nvidia-cuda-nvrtc-cu12==12.1.105 in /usr/local/lib/python3.10/dist-packages (from torch>=1.8.0->ultralytics) (12.1.105)
Requirement already satisfied: nvidia-cuda-runtime-cu12==12.1.105 in /usr/local/lib/python3.10/dist-packages (from torch>=1.8.0->ultralytics) (12.1.105)
Requirement already satisfied: nvidia-cuda-cupti-cu12==12.1.105 in /usr/local/lib/python3.10/dist-packages (from torch>=1.8.0->ultralytics) (12.1.105)
Requirement already satisfied: nvidia-cudnn-cu12==8.9.2.26 in /usr/local/lib/python3.10/dist-packages (from torch>=1.8.0->ultralytics) (8.9.2.26)
Requirement already satisfied: nvidia-cublas-cu12==12.1.3.1 in /usr/local/lib/python3.10/dist-packages (from torch>=1.8.0->ultralytics) (12.1.3.1)
Requirement already satisfied: nvidia-cufft-cu12==11.0.2.54 in /usr/local/lib/python3.10/dist-packages (from torch>=1.8.0->ultralytics) (11.0.2.54)
Requirement already satisfied: nvidia-curand-cu12==10.3.2.106 in /usr/local/lib/python3.10/dist-packages (from torch>=1.8.0->ultralytics) (10.3.2.106)
Requirement already satisfied: nvidia-cusolver-cu12==11.4.5.107 in /usr/local/lib/python3.10/dist-packages (from torch>=1.8.0->ultralytics) (11.4.5.107)
Requirement already satisfied: nvidia-cusparse-cu12==12.1.0.106 in /usr/local/lib/python3.10/dist-packages (from torch>=1.8.0->ultralytics) (12.1.0.106)
Requirement already satisfied: nvidia-nccl-cu12==2.19.3 in /usr/local/lib/python3.10/dist-packages (from torch>=1.8.0->ultralytics) (2.19.3)
Requirement already satisfied: nvidia-nvtx-cu12==12.1.105 in /usr/local/lib/python3.10/dist-packages (from torch>=1.8.0->ultralytics) (12.1.105)
Requirement already satisfied: triton==2.2.0 in /usr/local/lib/python3.10/dist-packages (from torch>=1.8.0->ultralytics) (2.2.0)
Requirement already satisfied: nvjitlink-cu12 in /usr/local/lib/python3.10/dist-packages (from nvidia-cusolver-cu12==11.4.5.107->torch>=1.8.0->ultralytics) (2.2.0)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.7->matplotlib>=3.3.0->ultralytics) (1.16.0)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from jinja2->torch>=1.8.0->ultralytics) (2.1.5)
Requirement already satisfied: mpmath>=0.19 in /usr/local/lib/python3.10/dist-packages (from sympy->torch>=1.8.0->ultralytics) (1.3.0)
```

```
import supervision as sv
import cv2
import numpy as np
import torch

from ultralytics import YOLO

model = YOLO('yolov8s.pt')
```

```
from google.colab import drive
drive.mount('/content/drive/')

Drive already mounted at /content/drive/; to attempt to forcibly remount, call drive.mount("/content/drive/", force_remount=True).
```

```

# 1 способ подсчета времени (смотрим разницу в прохождении одного отрезка соседними людьми и умножаем на кол-во людей в очереди)
video_info = sv.VideoInfo.from_video_path('/content/drive/MyDrive/data_sir/video/file7.mp4') # Сириус ближе к каске
byte_track = sv.ByteTrack(frame_rate=video_info.fps)
frames_generator = sv.get_video_frames_generator('/content/drive/MyDrive/data_sir/video/file7.mp4')
time = 0
polygon_entry = np.array([
[495, 315],[639, 403],[755, 319],[555, 215],[495, 311]
])
polygon_queue = np.array([
[8, 83],[8, 251],[192, 203],[808, 571],[1016, 567],[1028, 295],[244, 87],[100, 127],[101, 63],[7, 83]
])
polygon_table = np.array([
[39, 209],[55, 253],[175, 229],[159, 169],[51, 181],[35, 205]
])
polygon_check = np.array([
[206, 106],[218, 198],[182, 206],[170, 114],[206, 106]
])

box_annotator = sv.BoxAnnotator(thickness=1, text_thickness=1, text_scale=0.5)
zone_entry = sv.PolygonZone(polygon=polygon_entry, frame_resolution_wh=video_info.resolution_wh)
zone_queue = sv.PolygonZone(polygon=polygon_queue, frame_resolution_wh=video_info.resolution_wh)
zone_table = sv.PolygonZone(polygon=polygon_table, frame_resolution_wh=video_info.resolution_wh)
zone_check = sv.PolygonZone(polygon=polygon_check, frame_resolution_wh=video_info.resolution_wh)
zone_annotator_queue = sv.PolygonZoneAnnotator(zone=zone_queue, color=sv.Color.white(), thickness=3, text_thickness=3, text_scale=2)
zone_annotator_table = sv.PolygonZoneAnnotator(zone=zone_table, color=sv.Color.white(), thickness=3, text_thickness=3, text_scale=2)
zone_annotator_check = sv.PolygonZoneAnnotator(zone=zone_check, color=sv.Color.white(), thickness=3, text_thickness=3, text_scale=2)

entry_ids = []
last_entry_time = 0
times = []
with sv.VideoSink(target_path='/content/drive/MyDrive/data_sir/video/method_Ilya_two.mp4', video_info=video_info) as sink:
    for frame in frames_generator:
        time += 1/10
        results = model(frame, imgsz=1280)[0]
        detections = sv.Detections.from_ultralytics(results)
        detections = detections[detections.class_id==0]
        detections = byte_track.update_with_detections(detections=detections)
        zone_entry.trigger(detections=detections)
        CONST = 7 if (zone_check.trigger(detections=detections)==True).sum() > 0 else 0
        count = (zone_queue.trigger(detections=detections)==True).sum() - (zone_table.trigger(detections=detections)==True).sum() + CONST
        detections = detections[zone_entry.trigger(detections=detections)==True]
        print(detections.tracker_id)
        for id in detections.tracker_id:
            if id not in entry_ids:
                if len(entry_ids) == 0:
                    last_entry_time = time
                    times.append((time - last_entry_time) * count)
                    last_entry_time = time
                    entry_ids.append(id)
            else:
                print(f'people: {count}, id: {id}, time: {time}, last time: {last_entry_time}')
                times.append((time - last_entry_time) * count)
                last_entry_time = time
                entry_ids.append(id)

        box_annotator = sv.BoxAnnotator(thickness=1, text_thickness=1, text_scale=0.5)
        labels = [
            f'#{tracker_id}'
            for tracker_id
            in detections.tracker_id]
        frame = box_annotator.annotate(scene=frame, detections=detections, labels=labels)
        frame = zone_annotator_check.annotate(zone_annotator_queue.annotate(zone_annotator_table.annotate(scene=frame)))

        sink.write_frame(frame=frame)

print(f'average time for waiting: {sum(times)/len(times)}s')

```

▼ Текст заголовка по умолчанию

```

# @title Текст заголовка по умолчанию
# 2 way of calculating waiting time

# do everything the same except for other polygons
video_info = sv.VideoInfo.from_video_path('/content/drive/MyDrive/data_sir/video/Itog.mp4')
byte_track = sv.ByteTrack(frame_rate=video_info.fps)
frames_generator = sv.get_video_frames_generator('/content/drive/MyDrive/data_sir/video/Itog.mp4')
time = 0
polygon_left = np.array([
[11, 583],[19, 699],[171, 703],[159, 575],[15, 579]
])
polygon_right = np.array([
[964, 541],[1068, 489],[972, 393],[1068, 345],[1152, 405],[1132, 705],[980, 709],[964, 537]
])

box_annotator = sv.BoxAnnotator(thickness=1, text_thickness=1, text_scale=0.5)
zone_left = sv.PolygonZone(polygon=polygon_left, frame_resolution_wh=video_info.resolution_wh)
zone_right = sv.PolygonZone(polygon=polygon_right, frame_resolution_wh=video_info.resolution_wh)
zone_annotator_left = sv.PolygonZoneAnnotator(zone=zone_left, color=sv.Color.white(), thickness=3, text_thickness=3, text_scale=2)
zone_annotator_right = sv.PolygonZoneAnnotator(zone=zone_right, color=sv.Color.white(), thickness=3, text_thickness=3, text_scale=2)

# here we initialize twice more arrays (as we want to see the number of people enetring and exiting the queue)
entry_ids = []
exit_ids = []
entry_times = []
exit_times = []
with sv.VideoSink(target_path='queues_results_two.mp4', video_info=video_info) as sink:
    for frame in frames_generator:
        time += 1/video_info.fps
        results = model(frame, imgsz=1280)[0]
        detections = sv.Detections.from_ultralytics(results)
        detections = detections[detections.class_id==0]
        zone_left.trigger(detections=detections)
        zone_right.trigger(detections=detections)
        detections = byte_track.update_with_detections(detections=detections)

        for id in detections.tracker_id[zone_left.trigger(detections=detections)==True]:
            if id not in exit_ids:
                exit_times.append(time)
                exit_ids.append(id)
        for id in detections.tracker_id[zone_right.trigger(detections=detections)==True]:
            if id not in entry_ids:
                entry_times.append(time)
                entry_ids.append(id)

        if len(exit_times) >= 3:
            waiting_time = int(np.mean(np.asarray(exit_times[-3:]) - np.asarray(entry_times[len(exit_times)-3:len(exit_times)])))
            print(f'waiting time is {waiting_time}s')

        box_annotator = sv.BoxAnnotator(thickness=1, text_thickness=1, text_scale=0.5)

        labels_left = [
            f'#{tracker_id}'
            for tracker_id
            in detections.tracker_id[zone_left.trigger(detections=detections)==True]]
        labels_right = [
            f'#{tracker_id}'
            for tracker_id
            in detections.tracker_id[zone_right.trigger(detections=detections)==True]]

        frame = box_annotator.annotate(scene=frame, detections=detections[zone_left.trigger(detections=detections)==True], labels=labels_left)
        frame = box_annotator.annotate(scene=frame, detections=detections[zone_right.trigger(detections=detections)==True], labels=labels_right)
        frame = zone_annotator_left.annotate(zone_annotator_right.annotate(scene=frame))

        sink.write_frame(frame=frame)

```

```
0: 736x1280 31 persons, 1 bottle, 1 bowl, 11 chairs, 2 dining tables, 24.4ms
Speed: 5.6ms preprocess, 24.1ms inference, 2.0ms postprocess per image at shape (1, 3, 736, 1280)
waiting time is 86s
```

```
0: 736x1280 31 persons, 1 bottle, 1 bowl, 12 chairs, 2 dining tables, 24.4ms
Speed: 4.3ms preprocess, 24.4ms inference, 5.1ms postprocess per image at shape (1, 3, 736, 1280)
waiting time is 86s
```

```
SupervisionWarnings: BoxAnnotator is deprecated: `BoxAnnotator` is deprecated and will be removed in `supervision-0.22.0`. Use `
SupervisionWarnings: BoxAnnotator is deprecated: `BoxAnnotator` is deprecated and will be removed in `supervision-0.22.0`. Use `
SupervisionWarnings: BoxAnnotator is deprecated: `BoxAnnotator` is deprecated and will be removed in `supervision-0.22.0`. Use `
0: 736x1280 29 persons, 1 bottle, 1 bowl, 12 chairs, 2 dining tables, 25.0ms
Speed: 4.2ms preprocess, 25.0ms inference, 4.1ms postprocess per image at shape (1, 3, 736, 1280)
waiting time is 86s
```

```
0: 736x1280 24 persons, 1 bench, 1 handbag, 1 bottle, 1 bowl, 13 chairs, 2 dining tables, 24.9ms
Speed: 4.3ms preprocess, 24.9ms inference, 2.2ms postprocess per image at shape (1, 3, 736, 1280)
waiting time is 86s
```

```
0: 736x1280 26 persons, 1 bench, 1 bottle, 1 bowl, 13 chairs, 3 dining tables, 24.8ms
Speed: 4.2ms preprocess, 24.8ms inference, 3.2ms postprocess per image at shape (1, 3, 736, 1280)
waiting time is 86s
```

```
SupervisionWarnings: BoxAnnotator is deprecated: `BoxAnnotator` is deprecated and will be removed in `supervision-0.22.0`. Use `
SupervisionWarnings: BoxAnnotator is deprecated: `BoxAnnotator` is deprecated and will be removed in `supervision-0.22.0`. Use `
SupervisionWarnings: BoxAnnotator is deprecated: `BoxAnnotator` is deprecated and will be removed in `supervision-0.22.0`. Use `
0: 736x1280 26 persons, 1 bench, 1 bottle, 1 bowl, 10 chairs, 2 dining tables, 25.5ms
Speed: 3.7ms preprocess, 25.5ms inference, 4.5ms postprocess per image at shape (1, 3, 736, 1280)
waiting time is 86s
```

```
0: 736x1280 26 persons, 1 bench, 1 bottle, 1 bowl, 11 chairs, 2 dining tables, 24.8ms
Speed: 3.9ms preprocess, 24.8ms inference, 1.7ms postprocess per image at shape (1, 3, 736, 1280)
waiting time is 86s
```

```
0: 736x1280 28 persons, 1 bench, 1 bottle, 1 bowl, 13 chairs, 2 dining tables, 24.7ms
Speed: 3.3ms preprocess, 24.7ms inference, 1.7ms postprocess per image at shape (1, 3, 736, 1280)
waiting time is 86s
```

```
SupervisionWarnings: BoxAnnotator is deprecated: `BoxAnnotator` is deprecated and will be removed in `supervision-0.22.0`. Use `
SupervisionWarnings: BoxAnnotator is deprecated: `BoxAnnotator` is deprecated and will be removed in `supervision-0.22.0`. Use `
SupervisionWarnings: BoxAnnotator is deprecated: `BoxAnnotator` is deprecated and will be removed in `supervision-0.22.0`. Use `
```

```
import time
```

```
# подсчет кол-ва людей за столами (на нашей столовке)
start = time.time()
```

```
video_info = sv.VideoInfo.from_video_path('/content/drive/MyDrive/data_sir/video/table_itog.mp4')
polygon = np.array([
[554, 697],[62, 201],[62, 145],[474, 33],[1210, 237],[1210, 689],[558, 693]
])
zone = sv.PolygonZone(polygon=polygon, frame_resolution_wh=video_info.resolution_wh)
```

```
box_annotator = sv.BoxAnnotator(thickness=1, text_thickness=1, text_scale=0.5)
zone_annotator = sv.PolygonZoneAnnotator(zone=zone, color=sv.Color.white(), thickness=3, text_thickness=3, text_scale=2)
```

```
with sv.VideoSink(target_path='/content/drive/MyDrive/table_result.mp4', video_info=video_info) as sink:
```

```
cnt = 0
for frame in sv.get_video_frames_generator('/content/drive/MyDrive/data_sir/video/table_itog.mp4'):
    cnt += 1
    if cnt%1==0:

        results = model(frame, imgsz=1280)[0]
        detections = sv.Detections.from_ultralytics(results)
        detections = detections[detections.class_id == 0]
        zone.trigger(detections=detections)
        print(f'People at the tables {(zone.trigger(detections=detections))==True}.sum()')

        box_annotator = sv.BoxAnnotator(thickness=1, text_thickness=1, text_scale=0.5)
        labels = [f"{model.names[class_id]} {confidence:0.2f}" for confidence, class_id in \
            zip(detections.confidence, detections.class_id)]
        frame = box_annotator.annotate(scene=frame, detections=detections, labels=labels)
        frame = zone_annotator.annotate(scene=frame)

        sink.write_frame(frame=frame)

end = time.time()
print("The time of execution is :", (end-start), "s")
```

0: 736x1280 19 persons, 2 cups, 1 bowl, 32 chairs, 4 dining tables, 1 cell phone, 19.4ms  
Speed: 3.5ms preprocess, 19.4ms inference, 1.8ms postprocess per image at shape (1, 3, 736, 1280)  
People at the tables 18

0: 736x1280 19 persons, 3 cups, 1 bowl, 32 chairs, 4 dining tables, 1 laptop, 1 cell phone, 20.3ms  
Speed: 3.8ms preprocess, 20.3ms inference, 2.2ms postprocess per image at shape (1, 3, 736, 1280)  
People at the tables 18

SupervisionWarnings: BoxAnnotator is deprecated: `BoxAnnotator` is deprecated and will be removed in `supervision-0.22.0`. Use ``  
SupervisionWarnings: BoxAnnotator is deprecated: `BoxAnnotator` is deprecated and will be removed in `supervision-0.22.0`. Use ``  
SupervisionWarnings: BoxAnnotator is deprecated: `BoxAnnotator` is deprecated and will be removed in `supervision-0.22.0`. Use ``  
SupervisionWarnings: BoxAnnotator is deprecated: `BoxAnnotator` is deprecated and will be removed in `supervision-0.22.0`. Use ``  
0: 736x1280 19 persons, 3 cups, 1 bowl, 32 chairs, 4 dining tables, 1 cell phone, 20.5ms  
Speed: 3.8ms preprocess, 20.5ms inference, 1.8ms postprocess per image at shape (1, 3, 736, 1280)  
People at the tables 18

0: 736x1280 18 persons, 2 cups, 1 bowl, 32 chairs, 4 dining tables, 1 cell phone, 20.1ms  
Speed: 5.6ms preprocess, 20.1ms inference, 1.6ms postprocess per image at shape (1, 3, 736, 1280)  
People at the tables 18

0: 736x1280 19 persons, 2 cups, 1 bowl, 34 chairs, 3 dining tables, 1 cell phone, 20.1ms  
Speed: 5.9ms preprocess, 20.1ms inference, 2.5ms postprocess per image at shape (1, 3, 736, 1280)  
People at the tables 19

0: 736x1280 20 persons, 2 cups, 1 bowl, 33 chairs, 3 dining tables, 1 cell phone, 20.0ms  
Speed: 4.7ms preprocess, 20.0ms inference, 1.8ms postprocess per image at shape (1, 3, 736, 1280)  
People at the tables 19

SupervisionWarnings: BoxAnnotator is deprecated: `BoxAnnotator` is deprecated and will be removed in `supervision-0.22.0`. Use ``  
SupervisionWarnings: BoxAnnotator is deprecated: `BoxAnnotator` is deprecated and will be removed in `supervision-0.22.0`. Use ``  
SupervisionWarnings: BoxAnnotator is deprecated: `BoxAnnotator` is deprecated and will be removed in `supervision-0.22.0`. Use ``  
SupervisionWarnings: BoxAnnotator is deprecated: `BoxAnnotator` is deprecated and will be removed in `supervision-0.22.0`. Use ``  
0: 736x1280 19 persons, 1 handbag, 2 cups, 1 bowl, 34 chairs, 3 dining tables, 1 cell phone, 19.9ms  
Speed: 5.7ms preprocess, 19.9ms inference, 3.6ms postprocess per image at shape (1, 3, 736, 1280)  
People at the tables 19

0: 736x1280 19 persons, 1 handbag, 2 cups, 1 bowl, 34 chairs, 3 dining tables, 1 cell phone, 21.7ms  
Speed: 3.6ms preprocess, 21.7ms inference, 1.7ms postprocess per image at shape (1, 3, 736, 1280)  
People at the tables 19

0: 736x1280 19 persons, 2 cups, 1 bowl, 32 chairs, 3 dining tables, 1 cell phone, 19.6ms  
Speed: 4.8ms preprocess, 19.6ms inference, 1.7ms postprocess per image at shape (1, 3, 736, 1280)  
People at the tables 19

0: 736x1280 20 persons, 1 handbag, 2 cups, 1 bowl, 32 chairs, 3 dining tables, 1 cell phone, 19.6ms  
Speed: 3.5ms preprocess, 19.6ms inference, 2.3ms postprocess per image at shape (1, 3, 736, 1280)